

# PROXIMAL POLICY OPTIMIZATION (PPO)

Presenter: Majid Ghasemi

ECE 750 – October 2024

# AGENDA

- Policy Gradient Methods
- TRPO
- PPO
- How to do PPO?
- Applications
- Conclusion

# Policy Gradient (REINFORCE)

---

**Algorithm 1** REINFORCE Algorithm

---

```
1: REINFORCE( $s_0, \pi_\theta$ )
2: Initialize  $\pi_\theta$  to anything
3: while forever (for each episode) do
4:   Generate episode  $s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_T, a_T, r_T$  with  $\pi_\theta$ 
5:   for each step of the episode  $n = 0, 1, \dots, T$  do
6:      $G_n \leftarrow \sum_{t=0}^{T-n} \gamma^t r_{n+t}$ 
7:     Update policy:  $\theta \leftarrow \theta + \alpha \gamma^n G_n \nabla \log \pi_\theta(a_n | s_n)$ 
8:   end for
9: end while
10: return  $\pi_\theta$ 
```

---

update on  
each batch  
(trajectory)

$$g = \nabla_{\theta} J(\pi_{\theta}) = E_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{\infty} \gamma^t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \underbrace{A^{\pi_{\theta}}(s_t, a_t)} \right]$$

# Problems?

## Unstable Update

- Step size ( $\alpha$ )
  - If  $\alpha$  is too large  $\rightarrow$  bad policy
  - Generated batch from a bad policy  $\rightarrow$  bad samples are collected
  - Bad samples  $\rightarrow$  worse policy
- 
- $\alpha$  is too small?  $\rightarrow$  slow learning process

# Problems?

## Data Inefficiency


- On-policy method
- The data is thrown out after just one gradient update
- Complex neural networks -> training becomes very slow

# Data Inefficiency

- How to make it efficient?
- Avoid sampling from current policy?
- Use previous samples like replay buffer in DQN?

Current  
Policy

New  
Policy



X

Replay Buffer

old data



- Can we estimate an expectation of one distribution without taking samples from it?

**YES!**

$$E_{x \sim p}[f(x)] \approx \frac{1}{N} \sum_{i=1}^N f(x^i), \quad x^i \sim p$$

$$E_{x \sim p}[f(x)] = \int f(x) p(x) \, dx$$

$$= \int f(x) \frac{p(x)}{q(x)} q(x) \, dx$$

$$= E_{x \sim q} \left[ f(x) \frac{p(x)}{q(x)} \right]$$

# Importance sampling in Policy Gradient

$$= E_{x \sim q} \left[ f(x) \frac{p(x)}{q(x)} \right]$$

$$\nabla J(\theta) = E_{(s_t, a_t) \sim \pi_{\theta}^{\text{New}}} [\nabla \log \pi_{\theta}(a_t | s_t) A(s_t, a_t)]$$

$$= E_{(s_t, a_t) \sim \pi_{\theta_{\text{old}}}} \left[ \frac{\pi_{\theta}(s_t, a_t)}{\pi_{\theta_{\text{old}}}(s_t, a_t)} \nabla \log \pi_{\theta}(a_t | s_t) A(s_t, a_t) \right]$$

$$J(\theta) = E_{(s_t, a_t) \sim \pi_{\theta_{\text{old}}}} \left[ \frac{\pi_{\theta}(s_t, a_t)}{\pi_{\theta_{\text{old}}}(s_t, a_t)} A(s_t, a_t) \right]$$

→ surrogate objective function

# Problem?

- Two expectations are same, but we are using sampling method to estimate them -> variance

Close samples? Good to go

Far from each other? High variance

# Unstable update

- Confident updates  $\rightarrow$  New policy  $\sim$  Old policy
- Two approaches: Adaptive Learning Rate, Limit the policy update change

How can we measure the distance between two distributions?

# KL Divergence

KL Divergence is a tool to measure the distance of two distributions

# KL divergence of two policies

$$D_{\text{KL}}(\pi_1 || \pi_2)[s] = \sum_{a \in \mathcal{A}} \pi_1(a|s) \log \frac{\pi_1(a|s)}{\pi_2(a|s)}$$



# Trust Region Policy Optimization (TRPO)

$$\text{maximize}_{\theta} \hat{E}_t \left[ \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)} \hat{A}_t \right] \quad \text{subject to} \quad \hat{E}_t [\text{KL}[\pi_{\theta_{\text{old}}}(\cdot|s_t), \pi_{\theta}(\cdot|s_t)]] \leq \delta$$

$$\text{maximize}_{\theta} \hat{E}_t \left[ \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)} \hat{A}_t - \underline{\beta} \text{KL}[\pi_{\theta_{\text{old}}}(\cdot|s_t), \pi_{\theta}(\cdot|s_t)] \right]$$

TRPO uses a hard constraint rather than a penalty because it is hard to choose a single value of  $\beta$  that performs well across different problems—or even within a single problem, where the characteristics change over the course of learning.

# Proximal Policy Optimization (PPO)

- **TRPO uses conjugate gradient descent to handle the constraint**
- Hessian Matrix → expensive both in computation and space
- **Idea:**
- The constraint helps in the training process. However, maybe the constraint is not a strict constraint.
- Does it matter if we only break the constraint just a few times?
- What if we treat it as a "**soft**" constraint? Add proximal value to the objective function?

# PPO w/ Adaptive KL Penalty

$$L^{KL PEN}(\theta) = \hat{E}_t \left[ \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)} \hat{A}_t - \beta \text{KL}[\pi_{\theta_{old}}(\cdot|s_t), \pi_{\theta}(\cdot|s_t)] \right]$$

Hard to pick  $\beta$  value  $\rightarrow$  use adaptive  $\beta$

Compute  $d = \hat{E}_t[\text{KL}[\pi_{\theta_{old}}(\cdot|s_t), \pi_{\theta}(\cdot|s_t)]]$

- If  $d < d_{\text{targ}}/1.5$ ,  $\beta \leftarrow \beta/2$   $\rightarrow$  Relax
- If  $d > d_{\text{targ}} \times 1.5$ ,  $\beta \leftarrow \beta \times 2$   $\rightarrow$  Make Penalty

---

**Algorithm 2** PPO with Adaptive KL Penalty

---

- 1: **Input:** initial policy parameters  $\theta_0$ , initial KL penalty  $\beta_0$ , target KL-divergence  $\delta$
- 2: **for**  $k = 0, 1, 2, \dots$  **do**
- 3:     Collect set of partial trajectories  $\mathcal{D}_k$  on policy  $\pi_k = \pi(\theta_k)$
- 4:     Estimate advantages  $\hat{A}_t^{\pi_k}$  using any advantage estimation algorithm
- 5:     Compute policy update:

$$\theta_{k+1} = \arg \max_{\theta} \mathcal{L}_{\theta_k}(\theta) - \beta_k \overline{D}_{KL}(\theta || \theta_k)$$

- 6:     Take  $K$  steps of minibatch SGD (via Adam)
  - 7:     **if**  $\overline{D}_{KL}(\theta_{k+1} || \theta_k) \geq 1.5\delta$  **then**
  - 8:          $\beta_{k+1} \leftarrow 2\beta_k$
  - 9:     **else if**  $\overline{D}_{KL}(\theta_{k+1} || \theta_k) \leq \delta/1.5$  **then**
  - 10:          $\beta_{k+1} \leftarrow \beta_k/2$
  - 11:     **end if**
  - 12: **end for**
-

# PPO with Clipped Objective

$$\text{maximize}_{\theta} \hat{E}_t \left[ \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)} \hat{A}_t \right]$$

$$r_t(\theta) = \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)}$$

Invariance happens when  $r$  changes too quickly  $\rightarrow$  limit  $r$  within a range?

$$L^{CLIP}(\theta) = \hat{E}_t \left[ \min \left( r_t(\theta) \hat{A}_t, \underbrace{\text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)}_{\substack{\text{too small} \\ \text{too big}}} \hat{A}_t \right) \right]$$

---

**Algorithm 3** PPO with Clipped Objective

---

- 1: **Input:** initial policy parameters  $\theta_0$ , clipping threshold  $\epsilon$
- 2: **for**  $k = 0, 1, 2, \dots$  **do**
- 3:     Collect set of partial trajectories  $\mathcal{D}_k$  on policy  $\pi_k = \pi(\theta_k)$
- 4:     Estimate advantages  $\hat{A}_t^{\pi_k}$  using any advantage estimation algorithm
- 5:     Compute policy update

$$\theta_{k+1} = \arg \max_{\theta} L_{\theta_k}^{CLIP}(\theta)$$

by taking  $K$  steps of minibatch SGD (via Adam), where

$$L_{\theta_k}^{CLIP}(\theta) = E_{\tau \sim \pi_k} \left[ \sum_{t=0}^T \min \left( r_t(\theta) \hat{A}_t^{\pi_k}, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t^{\pi_k} \right) \right]$$

6: **end for**

---



# Practical Guide on PPO

## Components of PPO

- Policy Network
- Value Function

# Policy Network

State  $\rightarrow$  Neural Networks  $\rightarrow$  Probability of taking actions in that state

# Value Function

State  $\rightarrow$  Neural Networks  $\rightarrow$   $Q(s, a_1)$   
 $Q(s, a_2)$   
 $Q(s, a_3)$   
 $Q(s, a_4)$

# Applications?

- RLHF
- Discrete & Continuous states space
- You name some?